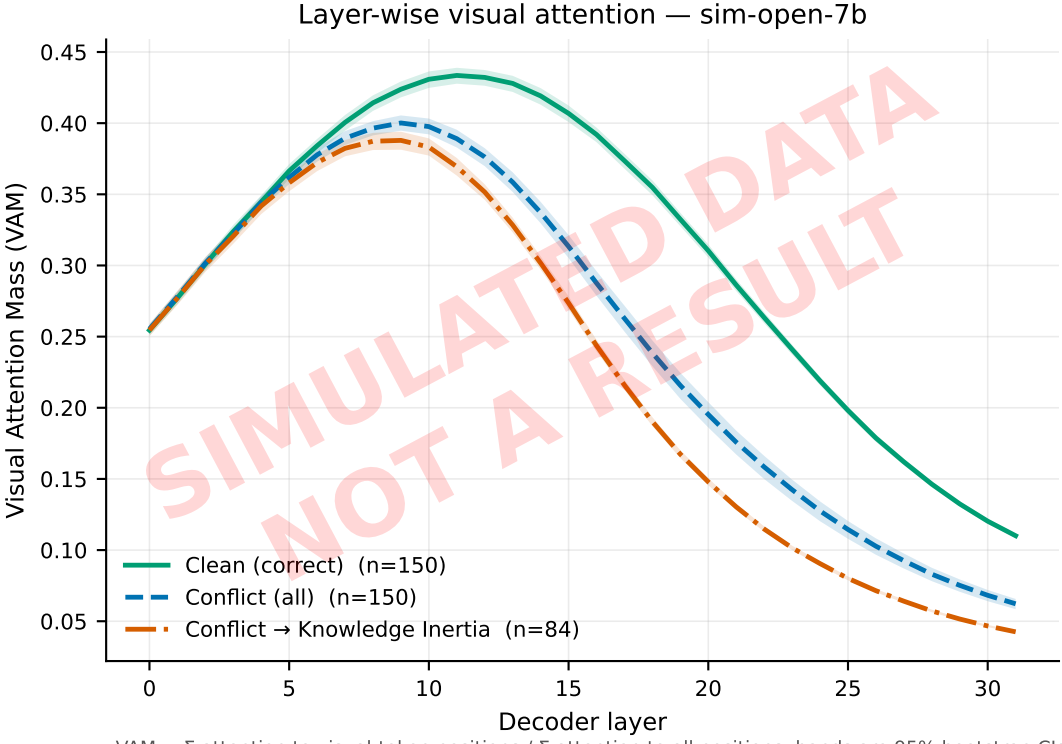
Layer-wise visual attention — sim-open-7b
Visual Attention Mass (VAM)
0.45
0.40
0.35
0.30
0.25
0.20
0.15
0.10
0.05
0
5
10
15
20
25
30
Decoder layer
Clean (correct) (n=150)
Conflict (all) (n=150)
Conflict → Knowledge Inertia (n=84)
SIMULATED DATA
NOT A RESULT
VAM = Σ attention to visual-token positions / Σ attention to all positions; bands are 95% bootstrap CIs